

Advanced Programming for Scientific Computing (PACS)

Lecture title: A brief introduction to floating
point numbers

Luca Formaggia

MOX
Dipartimento di Matematica
Politecnico di Milano

A.A. 2022/2023

Floating point number system

A normalized floating point number system F is formed by zero and real numbers of the form

$$y = \pm m \times \beta^{e-t}$$

where

β	basis (typically 2)	t	precision
m	mantissa, $\beta^{t-1} \leq m < \beta^t$	e	exponent, $e_{\min} \leq e \leq e_{\max}$

Or,

$$y = 0.d_1d_2\dots d_t \times \beta^e \quad \underline{d_1 \neq 0}, \quad 0 \leq d_i < \beta$$

we use all the space for the digits, we don't use extra precision

All floating point numbers representable in a computer belong to $F^* = F \cup F_s$ where F_s is the set of *subnormal numbers* of the form

$$y = \pm m \times \beta^{e_{\min}-t}, \quad 0 < m < \beta^{t-1}$$

The values of t , e_{\min} and e_{\max} characterize the **type** of floating point number and is defined by the **IEEE 754** standard.

Round off

A real number $x \in \text{range}(F^*)$ is approximated in a computer by $\hat{x} = fl(x) \in F^*$ and $e_r = (x - \hat{x})/x$ is the (relative) **rounding error**.
If $x \in \text{range}(F^*)$ then $\hat{x} = fl(x) \in F^*$ and $e_r = 0$. In general,

$$|e_r| \leq u = \frac{1}{2}\beta^{1-t}$$

where u is the **roundoff unit**. The machine epsilon ϵ_M is the smallest positive floating point number by which $fl(1 + \epsilon_M) \neq 1$.
We have:

$$u = \frac{1}{2}\epsilon_M$$

IEEE arithmetic

The IEEE 754 has been defined in 1985 (and amended various times since then) and defines the floating point arithmetic system normally implemented in modern processors. There are two main types of floating point numbers

Type	Size	t	e	u	Range
float	32	$23 + 1$	8	2^{-24}	$10^{\pm 38}$
double	64	$52 + 1$	11	2^{-53}	$10^{\pm 308}$

The value in the table indicate the number of bits used. Moreover, the implementation of IEEE arithmetic system should satisfy the **standard model**: if $x, y \in F$ then *we x and y are real (\mathbb{R}) numbers*

$$fl(x \text{ op } y) = (x \text{ op } y)(1 + \delta) \quad |\delta| \leq u, \quad \text{op} = + - \times /$$

\Rightarrow we have a control about the precision in the computation

Special numbers

The IEEE standard prescribes that

$$fl(x) = 0 \quad \text{if } |x| < F_{min} \quad (\text{UNDERFLOW})$$

$F_{min} \in F^*$ being the smallest positive floating point number.

$$fl(x) = \text{sign}(x)Inf \quad \text{if } |x| > F_{max} \quad (\text{OVERFLOW})$$

$F_{max} \in F$ being the maximum floating point number.

Moreover, $x/0 = \pm Inf$ if $x \neq 0$, $Inf + Inf = Inf$ and $x/Inf = 0$ if $x \neq 0$.

Finally, the special number NaN (**Not-a-Number**) indicates the result of an **invalid** operation ($0/0$, $\log(-1)$ etc) and we have $x \text{ op } NaN = NaN$, $Inf - Inf = NaN$ e $Inf/Inf = NaN$.

Floating point exceptions

The standard prescribes that in normal situations an invalid operation or an over/underflow **do not** stop computations, but produces the special numbers illustrated before.

However, the processor may record if an invalid floating point operation (normally called **floating point exception**) has occurred, so that the user may **trap** it.

We will discuss this issue into more detail later in the course, showing how you can capture floating point exceptions.

Forward and backward error

Let $f : \mathbb{R} \rightarrow \mathbb{R}$ and $\hat{f} : F \rightarrow F$ the corresponding expression on floating point numbers. Let $y = f(x)$ e $\hat{y} = \hat{f}(\hat{x})$. The analysis of the **forward error** aims to find a δ such that

$$e_f = \frac{|y - \hat{y}|}{|y|} \leq \delta$$

The relative *backward error* Δ is defined by $\hat{y} = f(x(1 + \Delta))$. If $f \in C^2$ we have that

$$\frac{|y - \hat{y}|}{|y|} = c(x)|\Delta| + O(\Delta^2)$$

with $c(x) = |xf'(x)/f(x)|$ called the relative **condition number** of f . In general, one looks for an estimate such that

$$e_f \leq C(x)|\Delta|$$

A simple example: cancellation

Let us consider $y = f(a, b) = a - b$ e $\hat{y} = f(a(1 + \Delta), b(1 + \Delta))$.
It is easy to find out that

$$\left| \frac{y - \hat{y}}{y} \right| = \frac{|a\Delta - b\Delta|}{|a - b|} \leq \frac{|a| + |b|}{|a - b|} |\Delta| \Rightarrow C = \frac{|a| + |b|}{|a - b|}$$

We have that $C \rightarrow \infty$ when $|a - b| \rightarrow 0$: the subtraction of almost equal floating point values causes a big round-off error. **We should avoid it whenever possible!**

A thorough analysis of floating point errors for common mathematical operations is found in the book by N.J. Higham **Accuracy and stability of numerical algorithms**, SIAM, ISBN: 978-0-89871-521-7, 2002.

An example: numerical differentiation

If $f(x) : \Omega \rightarrow \mathbb{R}$ is sufficiently regular we have

$$f'(x) = \frac{f(x+h) - f(x-h)}{2h} + O(h^2), \quad \forall x \in B_h(x).$$

Therefore the centered formula $\delta_f(x) = (f(x+h) - f(x-h))/2h$ is a second' order approximation of the derivative. However, the quantity actually computed is

$$\widehat{\delta}_f(x) = \frac{f[(x+h)(1+\Delta_1)] - f[(x-h)(1+\Delta_2)]}{2h},$$


with $|\Delta_1| \leq u$ and $|\Delta_2| \leq u$. If $f'(x) \neq 0$ we may obtain the estimate

$$|\widehat{\delta}_f(x) - \delta_f(x)| \leq \frac{C}{2h}$$

with $C \leq |x| |f'(x)| u$.

In the example [Examples/src/FloatingPoint/FinDiff](#) we also implement the fourth order formula $(\frac{1}{12}f(x-2h) - \frac{2}{3}f(x-h) - \frac{1}{12}f(x+2h) + \frac{2}{3}f(x+h))/h$ and we compute the numerical derivative of $100e^x$ at the point $x = 3$ with the two formulas and plot the error $|\hat{\delta}_f(x) - f'(x)|$ for different values of h , together with the estimated forward truncation error. We repeat the example using single, double and extended precision

The result for double precision



You may note that for $h \lesssim 10^{-6}$ the truncation error dominates the second order formula (Error double 2), while it already spoils the result for $h \lesssim 10^{-3}$ if we use the fourth order approximation!

Differentiation with(out) a difference


Sometimes one may use non-standard techniques if a great precision is needed. One may for instance note that if f is extended to a complex function and if the extension is analytic we may write (we need to apply Cauchy-Riemann equation and note that $f(x)$ is real)

$$f(\underbrace{x + ih}_{\text{complex number}}) = f(x) + ihf'(x) + O(h^2).$$

So $\text{Im}\left(\frac{f(x + ih)}{h}\right)$ is a second order approximation of $f'(x)$ and its computation does not imply taking differences!.

Indeed the result obtained for small h may be very precise. In [FloatingPoint/FinDiff](#) we have implemented also this formula. We show the result.

The result for double precision



You may note that with this trick (blue stars) we can reduce considerably the size of h . For larger h , as expected, we get results similar to the 2nd order formula.

Other nasty and less nasty examples

Some examples of floating point failures are in the [FloatingPoint](#) directory of the examples, with a file containing the explanation.

Some of them they are specially hand-crafted examples to highlight some possible unwanted side-effects of floating point operations.

Normally the situation is not that bad!

For instance, in [FloatingPoint/QuadraticRoot](#), where we show how the classic formula for the zero of a quadratic $ax^2 + bx + c$ may give incorrect results when $|b| \gg ac$, because of cancellation errors.

Finally, in [FloatingPoint/FPFailure](#) you find an example that shows some "floating point failures".

Numeric limits

C++ allows you to interrogate the characteristics of the numeric types [as implemented in your machine](#).

Not just floating points, but also integral types. In particular, we may look for the [machine epsilon](#) or the maximal representable number.

We need to use the Standard Library header `<limits>`:

```
#include <limits>
float eps=numeric_limits<float>::epsilon();
```

Since S++20 including the header `<numbers>` we have some mathematical constants in C++ style

See a complete example in [Examples/src/Numeric_Limits](#)

Beware of floating point comparisons!

Floating points have discrete values. So comparing them can be very critical. The statement

```
double a,b;  
... // many computations involving a and b  
if(a==b)....
```

is dangerous. Maybe the test is never satisfied. A stupid example

```
double a = std::pow(3.0,1./5);  
double b = a*a*a*a*a;  
double c =3.0;  
if(b==3.0) .. // IS FALSE!
```

It is better to write

```
if(std::abs(a-b)<tol)...
```

where `tol` is a well chosen tolerance (unfortunately it's not always evident what "well chosen" means.)